Assignment 7

1. **(Conjugate Priors)**

   (a) Consider the following form of the Normal distribution

   $$p(x \mid \mu, \kappa) = \frac{\kappa^{\frac{1}{2}}}{\sqrt{2\pi}} e^{-\frac{\kappa(x-\mu)^2}{2}}$$

   where κ (the variance inverse) is called the precision parameter. Show that this distribution can be written as an Exponential Family distribution of the form

   $$p(x \mid \theta_1, \theta_2) = h(x) e^{-\frac{\theta_1 x^2}{2} + \theta_2 x - \psi(\theta_1, \theta_2)}$$

   Characterize $h(x)$, (θ_1, θ_2) and the function $\psi(\theta_1, \theta_2)$.

   (b) Recall that the generic conjugate prior for an exponential family distribution is given by

   $$\pi(\theta_1, \theta_2) \propto e^{a_1 \theta_1 + a_2 \theta_2 - \gamma \psi(\theta_1, \theta_2)}. \tag{1}$$

   Substitute your expression for (θ_1, θ_2) from part (a) to show that the conjugate prior for the Normal model is of the form

   $$\pi(\kappa \mid a_0, b_0) \cdot \pi(\mu \mid \mu_0, \gamma \kappa) \propto \underbrace{\kappa^{a_0 - 1} e^{-\frac{\kappa}{b_0}}}_{\text{Gamma}(\kappa \mid a_0, b_0)} \cdot \underbrace{\kappa^{\frac{1}{2}} e^{-\frac{\gamma \kappa}{2}(\mu - \mu_0)^2}}_{\text{Normal}(\mu \mid \mu_0, \gamma \kappa)}.$$

   Your expressions for a_0, b_0 and μ_0 should be in terms of γ, a_1 and a_2. (This prior is known as the Normal-Gamma prior.)

   (c) Suppose $(\mu, \kappa) \sim$ Normal-Gamma$(a_0, b_0, \mu_0, \gamma)$, and the likelihood of the data, x, is $p(x \mid \mu, \kappa) = \frac{\kappa^{\frac{1}{2}}}{\sqrt{2\pi}} e^{-\frac{\kappa(x-\mu)^2}{2}}$. Compute the posterior distribution after you see N IID samples $\{x_1, \ldots, x_N\}$.

2. **(Order Restricted Inference)**
   Suppose one observes $y_1, \ldots, y_N$ where y_i is binomially distributed with sample size n_i and probability of success p_i, for $i = 1, \ldots, N$. The p_i's are unknown but domain specific knowledge tells us that

   $$0 \leq p_1 < p_2 < \cdots < p_N \leq 1. \tag{2}$$

   We therefore assume a uniform prior for $(p_1, \ldots, p_N)$ over the space in $\mathbb{R}^N$ defined by (2). Describe in detail a Gibbs sampling algorithm for sampling from the posterior distribution of $(p_1, \ldots, p_N)$ given $y_1, \ldots, y_N$.

3. **(Gibbs and the Hierarchical Normal Model)**
   Consider the hierarchical Normal model of Example 7 in the *MCMC and Bayesian Modeling* lecture notes. (This model is taken from Gelman et al's *Bayesian Data Analysis.*)

   (a) Write your own Gibbs sampler code in the language of your choice to sample from the posterior distribution.

   Hint: To simulate $X \sim$ Inv-$\chi^2(\nu, s^2)$ first simulate Y from the χ^2_ν distribution and then set $X = \nu s^2 / Y$.

   (b) Implement the Gelman-Rubin diagnostic by running 4 chains from over-dispersed starting points, discarding the first 50% of samples etc.

   (c) After running your code from (a) and (b) (and checking that the convergence diagnostics are satisfied!) report posterior quantiles (at the 2.5%, 25%, 50%, 75% and 97.5% levels) for θ_1, θ_2, θ_3, θ_4, μ, σ and τ. (Figure 1 displays results from Gelman et al's *Bayesian Data Analsyis*. You should obtain similar results.)

Estimand	\multicolumn{5}{c}{Posterior quantiles}	$\widehat{R}$
	2.5%	25%	median	75%	97.5%	
θ_1	58.9	60.6	61.3	62.1	63.5	1.01
θ_2	63.9	65.3	65.9	66.6	67.7	1.01
θ_3	66.0	67.1	67.8	68.5	69.5	1.01
θ_4	59.5	60.6	61.1	61.7	62.8	1.01
μ	56.9	62.2	63.9	65.5	73.4	1.04
σ	1.8	2.2	2.4	2.6	3.3	1.00
τ	2.1	3.6	4.9	7.6	26.6	1.05
$\log p(\mu, \log \sigma, \log \tau \| y)$	−67.6	−64.3	−63.4	−62.6	−62.0	1.02
$\log p(\theta, \mu, \log \sigma, \log \tau \| y)$	−70.6	−66.5	−65.1	−64.0	−62.4	1.01

Figure 1: Results for Exercise 3 from Gelman et al.'s *Bayesian Data Analysis.*

4. **(Exchangeable random variables)**
   We say $\mathbf{X} = (X_1, \ldots, X_N)$ is a vector of *exchangeable* random variables if there exists θ and a prior PDF $\pi(\theta)$ such that

$$\mathbb{P}(\mathbf{X} = \mathbf{x}) = \int \prod_{j=1}^{N} f(x_j \mid \theta)\pi(\theta)d\theta. \tag{3}$$

   It therefore follows from (3) that $X_1, \ldots, X_N$ are IID with density $f(\cdot \mid \theta)$ *given* θ.

(a) Show that

$$\mathbb{P}(X_k = x_k \mid \mathbf{X}_{-k} = \mathbf{x}_{-k}) = \int f(x_k \mid \theta) p(\theta \mid \mathbf{X}_{-k} = \mathbf{x}_{-k}) d\theta$$

where $p(\theta \mid \mathbf{X}_{-k} = \mathbf{x}_{-k})$ denotes the posterior density of θ given $\mathbf{X}_{-k}$.

(b) Consider the following special case where under the prior distribution $\theta \sim \text{Beta}(\alpha, 0)$, and $f(x \mid \theta) = \text{Bernoulli}(\theta)$. Show that

$$\mathbb{P}(X_k = 1 \mid \mathbf{X}_{-k}) = \frac{\alpha + m_{-k}}{\alpha + N - 1}$$

where $m_{-k} = \sum_{j \neq k} \mathbf{1}(X_j = 1)$.

(c) Continuing on from part (b), suppose the X_k's are not observable. Instead for each X_k we observe a variable Y_k that is distributed according to the conditional distribution $g(Y \mid X)$. Let $\mathbf{Y} = (Y_1, \dots, Y_N)$ denote the observed values of Y. Show that

$$\mathbb{P}(X_k = 1 \mid \mathbf{X}_{-k}, \mathbf{Y}) \propto g(Y_k \mid X_k = 1) \cdot \frac{\alpha + m_{-k}}{\alpha + N - 1}.$$

Remark: Note that the results of this question provide all the conditional distributions that you would need for a Gibbs sampler in this important class of models.

5. **(Optional! Convergence Diagnostics)**
In the lecture slides we defined

$$\widehat{\text{Var}}^+ (\psi \mid \mathbf{X}) := \frac{n-1}{n} W + \frac{1}{n} B \tag{4}$$

where

$$B := \frac{n}{m-1} \sum_{j=1}^{m} \left(\bar{\psi}_{\cdot j} - \bar{\psi}_{\cdot \cdot} \right)^2$$

$$W := \frac{1}{m} \sum_{j=1}^{m} s_j^2 \quad \text{where} \quad s_j^2 := \frac{1}{n-1} \sum_{i=1}^{n} \left(\psi_{ij} - \bar{\psi}_{\cdot j} \right)^2.$$

These definitions were based on having m chains each with n samples after discarding the burn-in samples and ψ is some scalar function of the parameters / hidden variables over which the posterior is defined. We claimed that $\widehat{\text{Var}}^+ (\psi \mid \mathbf{X})$ was an unbiased estimator for $\text{Var}^+ (\psi \mid \mathbf{X})$ under stationarity. In this question, we will justify this claim.

(a) Suppose $Y_1, \ldots, Y_n$ is a sample from a stationary process with mean μ and auto-covariance function $\gamma(h)$. Show that

$$\mathrm{Var}(\bar{Y}) = \frac{\gamma(0)}{n} R_n \tag{5}$$

where $R_n := 1 + 2\sum_{h=1}^{n-1} \rho(h)\left(1 - \frac{h}{n}\right)$ and $\rho(h) := \gamma(h)/\gamma(0)$ is the autocorrelation function. Note that $\gamma(0) = \mathrm{Var}(Y)$. (If you don't know what the autocovariance function is try Google, Wikipedia or any time-series book.) Most stationary processes generated by MCMC have $\rho(h) \geq 0$ so that if we use (5) to estimate $\mathrm{Var}(Y)$ then we need to take this autocorrelation into account.

(b) Suppose now that Y follows an $AR(1)$ process (a reasonable approximation to an MCMC process) so that $Y_n = \phi Y_{n-1} + \epsilon$. In that case it is straightforward to check that $\rho(h) = \phi^h$. Now justify the approximation

$$R_n \approx \frac{1 + \phi}{1 - \phi}.$$

(c) Use the identity

$$\sum_{i=1}^{n}(Y_i - \mu)^2 = \sum_{i=1}^{n}(Y_i - \bar{Y})^2 + n(\bar{Y} - \mu)^2$$

and (5) to show that $\mathrm{E}\left[\sum_{i=1}^{n}(Y_i - \bar{Y})^2\right] = \gamma(0)(n - R_n)$. Argue then that

$$\widehat{\mathrm{Var}}(Y) := \frac{\sum_{i=1}^{n}(Y_i - \bar{Y})^2 + \widehat{\gamma(0)R_n}}{n}$$

is an unbiased estimator of $\mathrm{Var}(Y)$ when $\widehat{\gamma(0)R_n}$ is an unbiased estimator of $\gamma(0)R_n$.

(d) Explain how you could construct such an unbiased estimator of $\gamma(0)R_n$ using m realizations (each of length n) of the process. Now justify (4).